

# Cloud and Web Applications


## **Part 3 : Server-Side Techniques**

# Table of Contents

- J2EE
- Node.js
- Database System
  - RDBMS
  - MongoDB
- REST
- Develop Web Apps for Cloud
- Cloud Architecture Techniques

# **J2EE**

# J2EE Technology Overview



# J2EE Technology Overview (2)

- J2EE is a specification that different vendors implement
- J2EE Application Servers
  - For running full stack(Servlets, JSP, EJB)
  - jBoss, Oracle WebLogic, GlassFish, IBM Websphere ...
- Java WebServers
  - For running Servlets and JSP
  - Tomcat, Jetty,
- EJB-only server
  - OpenEJB

# J2EE Technology Specifications


- Java API for RESTful Web Services (JAX-RS)
- Web Services 1.3 [JSR109](#)
- Java API for XML-Based Web Services (JAX-WS) 2.2 [JSR224](#)
- Java Architecture for XML Binding (JAXB) 2.2 [JSR222](#)
- Web Services Metadata for the Java Platform 2.1 [JSR181](#)
- Java API for XML-based RPC (JAX-RPC) 1.1 [JSR101](#)
- Java APIs for XML Messaging (JAXM) 1.3 [JSR67](#)
- Java API for XML Registries (JAXR) 1.0 [JSR93](#)
- Java Servlet 3.0 [JSR315](#)
- JavaServer Faces (JSF) 2.0 [JSR314](#)
- JavaServer Pages (JSP) 2.2 [JSR245](#)
- Expression Language (EL) 2.2 [JSR245](#)
- JavaServer Pages Standard Tag Library (JSTL) 1.2 [JSR52](#)
- Enterprise JavaBeans (EJB) 3.1 [JSR318](#) Lite
- Java Persistence API (JPA) 2.0 [JSR317](#)
- Contexts and Dependency Injection for Java 1.0 [JSR299](#)
- Dependency Injection for Java 1.0 [JSR330](#)
- Common Annotations for the Java Platform 1.1 [JSR250](#)
- Java Message Service API (JMS) 1.1 [JSR914](#)
- Java Transaction API (JTA) 1.1 [JSR907](#)
- JavaMail API 1.4 [JSR919](#)
- Java Authentication Service Provider Interface for Containers (JASPIC) 1.0 [JSR196](#)
- Java Authorization Service Provider Contract for Containers (JACC) 1.4 [JSR115](#)
- ...

# Web Technologies in J2EE


- Java Servlet 3.0
- JavaServer Faces 2.0
- JavaServer Pages 2.2/Expression Language 2.2
- Standard Tag Library for JavaServer Pages (JSTL)
- Debugging Support for Other Languages 1.0

# Java Servlets

- A  **servlet** is a Java program that is invoked by a web server in response to a request.
- Hosted by a  **servlet container**




The web server handles the HTTP transaction details



Example: Apache Tomcat, both a web server and a servlet container

# Servlet Operation



# Servlet Example

- This servlet will say "Hello!" (in HTML)

```
package servlet;
import javax.servlet.http.*;

public class HelloServlet extends HttpServlet {
    public void service(HttpServletRequest req,
        HttpServletResponse res) throws IOException {
        PrintWriter htmlOut = res.getWriter();
        res.setContentType("text/html");
        htmlOut.println("<html><head><title>" +
            "Servlet Example Output</title></head><body>" +
            "<p>Hello!</p>" + "</body></html>");
        htmlOut.close();
    }
}
```

# Servlet Configuration

- The web application configuration file, `web.xml`, identifies servlets and defines a mapping from requests to servlets

An identifying name for the servlet (appears twice)

```
<servlet>
  <servlet-name>HelloServlet</servlet-name>
  <servlet-class>servlet.HelloServlet</servlet-class>
</servlet>
<servlet-mapping>
  <servlet-name>HelloServlet</servlet-name>
  <url-pattern>/hello</url-pattern>
</servlet-mapping>
```

The pathname used to invoke the servlet  
(relative to the web application URL)


The servlet's package  
and class names

# Java Server Pages (JSP)

- HTML document with embedded Java code
- Translated into Java servlets when first used
- Allow static HTML content to be separated from dynamic content

```
<%@ page %>
<html>
<head><title>Example Page</title></head>
<body>
<h1>Example Page</h1>
<%
    String name = req.getParameter("name");
%>
<p>Hello, <%= name %></p>
...
</html>
```

# Execution of JSP




# Java Beans

- A Java Bean is a Java class that can be invoked from a JSP
- This allows complex processing required by a JSP to be separated from the JSP
- A JSP can instantiate an instance of a bean class and invoke its methods

# Java Bean Uses

- Java Beans are typically used to encapsulate complex application logic, including calculations, database access, or network access



# JSP and Java Bean Example

- Java Bean


```
package bean;
public class AdderUtility {
    public int add(int a, int b) {
        return a + b;
    }
}
```

- JSP

```
<%@ page %>
<jsp:useBean id="adder" class="bean.AdderUtility"
    scope="session" />

<%
    int a, b, sum;
    ...
    int sum = adder.add(a, b);
%>
```

# MVC in J2EE



1. User requests "login" option; Controller sets "login.jsp" as the next view.
2. User enters credentials; Controller invokes LoginHandler to handle transaction
3. etc.

# Traditional J2EE Users

- Examples of use
  - Insurance companies
  - Online banks
  - Manufacturing industries
  - Public sector

*'Cobol of the 21st century'*

# New J2EE Techniques

- Many languages with seamless integration: you can share libraries, code, classes, etc.
  - Groovy
    - Object-oriented
    - Inspired by Java, Python, Ruby, Smalltalk
  - Scala
    - Functional, object-oriented
    - "Cutting away Java's syntactic overhead, adding power"
    - Inspired by Java, Scheme, Erlang, Haskell, Lisp
  - JRuby: implementation of Ruby on JVM
  - Jython: implementation of Python on JVM

# Extensions to Core J2EE

- Grails (Groovy on Rails)
  - Interpreted script compiled to Java bytecode at runtime
  - Supports writing Java code directly
  - Java libraries, JPA models etc directly available
- Spring / Spring MVC
  - Abstraction, no need to work directly with Servlets etc
    - XML-based configuration
  - IoC/DI
  - Java EE 6 based on Spring
- JSF - Java Server Faces
  - Notation for generating JSF-pages (View), which communicate with managed beans(Controller)
- Struts
  - Custom tag libraries
  - XML-based

# J2EE Trends

- Moving from single-tier or two-tier to **multi-tier** architecture
- Moving from monolithic model to **object-based** application model
- Moving from application-based client to HTML-based client

# “Trendy” J2EE users: LinkedIn

- Started with Java platform, using Java EE and extensions
  - Spring Framework
  - Grails
- Now utilizing also Scala and JRuby
  - Scala for back-end processing
  - JRuby for integration interfaces

# “Trendy” J2EE users: Twitter

- Started with Ruby on Rails
- Now using Java and Scala in back-end processing
  - Why?
    - Scalability and Performance
    - SOA
    - Encapsulation (re-use, maintenance)

# “Trendy” J2EE users: Others

- Google, Amazon and many others use J2EE
- What about Facebook?
  - Writing PHP, which quickly lead to serious performance issues
  - Started compiling PHP to C++
  - Are now investigating using PHP on JVM

# **NODE.JS**

# Introduction to Node.js

“Node.js is a platform built on Chrome's JavaScript runtime for easily building fast, scalable network applications. Node.js uses an event-driven, non-blocking I/O model that makes it lightweight and efficient, perfect for data-intensive real-time applications that run across distributed devices.”

- nodejs.org



# Preface

- Developed by Ryan Dahl pada tahun 2009
- Server Side JavaScript
- Built on Google's V8
- An environment for developing high performance web services
- Using event-driven, asynchronous I/O to minimize overhead and maximize scalability.
- The goal is to provide an easy way to build scalable network servers

# Threads vs. Event-Driven

<b>Threads</b>	<b>Asynchronous Event-driven</b>
Lock application / request with listener-workers threads	only one thread, which repeatedly fetches an event
Using incoming-request model	Using queue and then processes it
multithreaded server might block the request which might involve multiple events	manually saves state and then goes on to process the next event
Using context switching	no contention and no context switches
Using multithreading environments where listener and workers threads are used frequently to take an incoming-request lock	Using asynchronous I/O facilities (callbacks, not poll/select or O_NONBLOCK) environments

# Threads vs. Event-Driven (2)



# Why Node.js Use Event-based?

In a normal process cycle the web server while processing the request will have to wait for the I/O operations and thus blocking the next request to be processed.

Node.JS process each request as events, The server doesn't wait for the IO operation to complete while it can handle other request at the same time.

When the I/O operation of first request is completed it will call-back the server to complete the request.

# Node.js Potential Wins

- JavaScript is the popular language of the web
- Web applications spend most of their time doing I/O
- Can implement the same programming language on client and server. Code can be migrated between server and client more easily
- Common data formats (JSON) between server and client
- Common software tools, testing or quality reporting tools for server and client
- Node.js ships with a lot of useful modules, so you don't have to write everything from scratch.
- Thus, Node.js is really two things: a runtime environment and a library

# Node.js vs. Apache

<b>Platform</b>	<b>Number of request per second</b>
PHP ( via Apache)	3187,27
Static ( via Apache )	2966,51
Node.js	5569,30



# Node.js vs. PHP + Nginx

<b>Platform</b>	<b>Static (number of request per second)</b>	<b>Query MySQL (number of request per second)</b>
PHP + Nginx	3624	1293
Node.js	7677	2999

# Node.js at PayPal

- PayPal's web applications are moving away from Java and onto JavaScript and Node.js.
- Comparing with Java, the Node.js app was:
  - Built almost **twice as fast with fewer people**
  - Written in **33% fewer lines of code**
  - Constructed with **40% fewer files**

# Performance of PayPal App.



- Double the requests per second vs. the Java application.
- 35% decrease in the average response time for the same page. This resulted in the pages being served 200ms faster.


# Lab : Open Command Prompt

Cmd is the command line interface tool used to execute the various commands.

- Win XP: **Start button** -> **Run** -> type **cmd.exe**, press enter.
- Win 7: **Start button** -> type **cmd** in the search box -> Click on **cmd** the search results listing.
- Win 8: Swipe to **Apps** screen -> locate *Windows System* -> click on **Command Prompt**
- Win 10: Press **win+R**, type **cmd.exe** in the prompt box, press enter.

# Basic cmd commands


- `dir` : display a list of files and folders
- `cd <folder>` : change directory - move to a specific *folder*



- `md <name>` : create a new folder with *<name>* in the current folder.
- `ping <ip>` : test a network connection to the ip


# Lab : Install NodeJS

- Download and install node.js : <http://nodejs.org/>
  - Click ‘Install’ or go to the ‘Downloads’ page
  - Once downloaded, run the installer
  - Install Node.js in directory **c:\>nodejs**



# Test the Installation

- Open a cmd, go to the **nodejs** directory
- Type **node -v** in command line
- If the version information is shown, Node.js is correctly installed.



```
命令提示符

C:\Users\lidan>cd c:\nodejs

c:\nodejs>node -v
v0.10.28


c:\nodejs>
```

# A Basic HTTP Server

- Create `helloworld.js` file in the nodejs folder as the follows


```
var http = require("http");
http.createServer(function(request, response) {
    response.writeHead(200, {"Content-Type": "text/html"});
    response.write("Hello World!");
    response.end();
}).listen(8080);
console.log("Server has started.");
```

- Execute the js file in Node.js:



# Test the HTTP Server

- Open your browser and access <http://localhost:8080>. This should display a web page that says “Hello World”.



# Analyzing the HTTP Server

```
var http = require("http");

http.createServer(function(request, response) {
    response.writeHead(200, {"Content-Type": "text/html"});
    response.write("Hello World!");
    response.end();
}).listen(8080);

console.log("Server has started.");
```

*requires the “http” module from Node.js library and assigns it to variable http*

*the parameter of the createServer : an anonymous function to callback when a http request is arrived*

*Invokes method listen of the object , listen to a http port*

*calls function createServer of the http module, returns an object*

# Another Version of the HTTP Server

```
var http = require("http") ;

function onRequest(request, response) {
    console.log("Request received.");

    response.writeHead(200, {"Content-Type": "text/plain"});
    response.write("Hello World");
    response.end();
}

http.createServer(onRequest).listen(8080);

console.log("Server has started.");
```

# **DATABASE SYSTEMS**

# RDBMS

- A Relational Database Management System (RDBMS) provides storage and access for structural data in web applications
- Examples: MySQL, SQL Server, Oracle, PostgreSQL

# Other Types of Databases

- Hierarchical database
  - organizes data in a tree-like structure
  - defines a parent/child relationship
  - each parent can have many children but each child has only one parent
  - folder or directory structure for storing files on a computer is an example of a hierarchical data structure
- NoSQL- promotes the use of non-relational databases and does not require fixed table schemas as with the relational model

# Relational Database Structure

- A database includes one or more **tables**
  - Each table represents one *type* of entity
- Example: Tables in a Library Database

User

Loan  
(transaction)

Recording

Book

# Relational Database Structure (2)

- Each table **field** represents an entity attribute
- Each table **row** represents one entity

Car table:

<b>Year</b>	<b>Make</b>	<b>Model</b>	<b>Color</b>
1973	Volkswagen	Jetta	Red
1992	Ford	Aerostar	Blue
2004	Chevrolet	Suburban	Black

row →

field ↗

# Structured Query Language (SQL)

- Development since 1970 by IBM
- A standard language for creating and maintaining relational databases
- SQL statement types:
  - *data definition*: create databases and tables
  - *data manipulation*: add, modify, delete data
  - *data control*: set access permissions

# Basic SQL Statements

- Data definition
  - CREATE, DROP
- Data manipulation
  - SELECT, INSERT, UPDATE, DELETE
- Data control
  - GRANT, REVOKE

# CREATE

- Create a database or a table

```
CREATE DATABASE mydata
```

```
CREATE mydata.player (
    playerNr int PRIMARY KEY,
    name VARCHAR(30),
    isCurrent BOOLEAN NOT NULL)
```

# Example table creation

```
create table student  
(  
    id int,  
    name varchar(255),  
    major char(4),  
    gender char(1),  
    dob date,  
    constraint student_pk primary key (id)  
);
```

why did we specify  
these attribute types?

student\_pk is an arbitrary name

# Basic SQL Data Types

- INTEGER
- DECIMAL(T, R)
  - T=total digits, R=right digits (after '.')
- FLOAT
- CHAR(N)            N characters
- VARCHAR(N)        up to N characters
- BOOLEAN
- DATE
- TIME

# DROP

- DROP can be used to delete an entire database or a table

```
DROP mydata
```

```
DROP mydata.player
```

# SELECT

- SELECT retrieves data from a database

```
SELECT field-list FROM database.table  
WHERE condition  
ORDER BY field-list
```

- *field-list* is a comma-separated list of fields from the named table (\* means "all fields")
- *condition* is a Boolean condition using field names and/or constants

# SELECT Example

- Example

```
SELECT * FROM mydata.player
```

```
SELECT playerNr, name FROM player  
WHERE isCurrent=TRUE
```

```
SELECT playerNr, name, status FROM player  
WHERE playerNr >= 90001  
ORDER BY status, name
```

# INSERT

- INSERT adds a new row to a table

```
INSERT INTO player  
VALUES (23752, 'Jane Doe', TRUE)
```

# UPDATE

- UPDATE changes one or more rows

```
UPDATE database.table  
    SET field-assignment-list  
    WHERE condition
```

```
UPDATE player  
    SET isCurrent=TRUE  
    WHERE playerNr=33256
```

# DELETE


- DELETE removes one or more rows from a table

```
DELETE FROM database.table  
WHERE condition
```

```
DELETE FROM player  
WHERE playerNr=33523
```


# JDBC

- Java Database Connectivity (JDBC) is a Java API that allows Java programs to interact with relational database management systems
- Interaction also requires a **database driver**, which translates JDBC commands to procedure calls on the RDBMS



# JDBC Design

- An effective design for database access:
  - JSP: user interface presentation
  - Java Servlet: application logic
  - Java Bean: database access (JDBC)



# NoSQL Databases

- A NoSQL database provides a mechanism for storage and retrieval of data that employs less constrained consistency models than traditional relational databases.
- NoSQL databases only support eventual consistency which is a consistency model used in distributed computing that informally guarantees that, if no new updates are made to a given data item, eventually all accesses to that item will return the last updated value.

# NoSQL Databases (2)

- NoSQL databases are often highly optimized key-value stores intended for simple retrieval and appending operations, with the goal being significant performance benefits in terms of latency and throughput.
- Key-value stores allow the application to store its data in a schema-less way.
  - The data could be stored in a datatype of a programming language or an object. Because of this, there is no need for a fixed data model.

# NoSQL Market Data

2014  
8%

% NoSQL Enterprise Adoption

2010  
5%    2015  
20%



*"MongoDB is the new MySQL."*


FORRESTER®

*"Current adoption of NoSQL in enterprises is 4% in 2010, expected to double by 2014, and grow to 20% by 2015. 10% of existing apps could move to NoSQL. [Sept 2011]"*

Gartner®

*"NoSQL is in its infancy, many immature with only community support and limited tool capability; however, expect to see a few used more widely during the next 5 years."*

- 2011: Market Research Media noted worldwide NoSQL market was expected to reach ~\$3.4B by 2018, generating \$14B from 2013-2018 with a CAGR of 21%
- Comparison: data implies NoSQL market ~\$895M
  - MySQL in 2011 ~\$100M



# RDBMS vs. NoSQL

- RDBMS
  - Transactional guarantee, complex functionality, Data consistency, correctness
  - Hard to scale
- NoSQL
  - Simple key-value data model, schema free, simple functionality, relaxed consistency
  - Horizontal scalability, high availability
- NewSQL = RDBMS + NoSQL



mongoDB

# MongoDB Overview


- Document oriented, not table/row oriented
- Collection of binary JSON (BSON) documents
- Schemaless
- No relations or transactions native in database
- Scalable and high-performance
- Fast In-Place Updates
- Map/Reduce
- Full index support
- Servers for all major platforms
- Drivers for all major development environments
- Free and open-source, but also commercial support

# Why MongoDB

- Rapid Application Prototyping
  - Schema rides as objects inside the JSON Script, not as a separate item.
    - Schemas change quite quickly in modern apps as business needs change.
    - Apps have shorter production lives and even shorter life-cycles as a result
- Scalability
  - Automatic data rebalancing across cheap, inexpensive commodity network hardware during the shard operation.
- Large amount of customers worldwide.
- Free basic edition
- Many related free software downloads

# Use cases

- High performance and scalable applications
- Most web applications where you would previously use SQL



Do not use for:

- Transaction-critical applications

# BSON - Data Type in MongoDB

- BSON is the binary representation of JSON.
- MongoDB drivers send and receive in BSON form.

```
{"hello":      → "\x16\x00\x00\x00\x02hello\x00  
"world"}          \x06\x00\x00\x00world\x00\x00"  
  
{ "BSON":      → "\x31\x00\x00\x00\x04BSON\x00\x26\x00  
["awesome",    \x00\x00\x020\x00\x08\x00\x00\x00  
5.05, 1986]}   → \x00awesome\x00\x011\x00\x33\x33\x33\x33\x33\x33  
                  \x14\x40\x102\x00\xc2\x07\x00\x00  
                  \x00\x00"
```

# Terminology and Concepts

RDBMS		MongoDB
Database	→	Database
Table, View	→	Collection
Row	→	Document (JSON, BSON)
Column	→	Field
Index	→	Index
Join	→	Embedded Document
Foreign Key	→	Reference
Partition	→	Shard

# Schema Free

- MongoDB does not need any pre-defined data schema
- Every document could have different data!

```
{name: "will",  
eyes: "blue",  
birthplace: "NY",  
aliases: ["bill", "la  
ciacco"],  
loc: [32.7, 63.4],  
boss: "ben"}
```

```
{name: "jeff",  
eyes: "blue",  
loc: [40.7, 73.4],  
boss: "ben"}
```

```
{name: "brendan",  
aliases: ["el diablo"]}
```

```
{name: "ben",  
hat: "yes"}
```

```
{name: "matt",  
pizza: "DiGiorno",  
height: 72,  
loc: [44.6, 71.3]}
```



# CRUD

- Create
  - `db.collection.insert( <document> )`
  - `db.collection.save( <document> )`
  - `db.collection.update( <query>, <update>, { upsert: true } )`
- Read
  - `db.collection.find( <query>, <projection> )`
  - `db.collection.findOne( <query>, <projection> )`
- Update
  - `db.collection.update( <query>, <update>, <options> )`
- Delete
  - `db.collection.remove( <query>, <justOne> )`

# Example-Querying

```
> db.users.findOne({age:39})  
> db.users.find({'last': 'Doe'})  
// retrieve all users order by last_name:  
> db.users.find({}).sort({last: 1});
```

```
-   { "_id" : ObjectId("5114e0bd42..."),  
    "first" : "John",  
    "last" : "Doe",  
    "age" : 39,  
    "interests" : [  
        "Reading",  
        "Mountain Biking"]  
    "favorites": {  
        "color": "Blue",  
        "sport": "Soccer"} }
```

# Advanced Querying

```
{ name: "Joe", address: { city: "San Francisco", state: "CA" }, likes: [ 'scuba', 'math', 'literature' ] }
```

// field in sub-document:

```
db.persons.find( { "address.state" : "CA" } )
```

// find in array:

```
db.persons.find( { likes : "math" } )
```

// regular expressions:

```
db.persons.find( { name : /acme.*corp/i } );
```

// javascript where clause:

```
db.persons.find("this.name != 'Joe'");
```

// check for existence of field:

```
db.persons.find( { address : { $exists : true } } );
```

# Insert & Update

- Supports bulk inserts
- Default saves are upserts
- In place updating
- Atomic transactions for single document
- Server side JavaScript execution

# Insert & Update Example

```
> db.user.insert({  
  first: "John",  
  last : "Doe",  
  age: 39  
})
```

```
> db.user.update(  
  {"_id" : ObjectId("51...")},  
  {  
    $set: {  
      age: 40,  
      salary: 7000}  
  })
```

```
> db.user.remove({  
  "first": /^J/  
})
```


# Indexes

- Unique index on primary key (`_id` field)
- Create index from application code (`ensureIndex`)
- Index on embedded documents and fields
- Index on array fields (*multikey index*)
- Geospatial index
- No native full text indexing

# Auto-sharding

- Partitions data across *shards*
- Any BSON document resides on only one shard
- Increases write capacity and total data size
- Data automatically distributed
- Sharding transparent to application layer
- Partitioning based on client-defined shard key
- Good shard keys are highly distributed in value and write operations
- Sharding requires config servers (minimal 3) to maintain metadata

# Mongo Sharding



# GridFS

- Store and retrieve files that exceed the BSON-document size limit of 16MB.
- Large blob data, limited only by storage space
- Supports many thousands of files
- Supports often changing files




# Lab : Install MongoDB

- Download mongoDB from  
<http://www.mongodb.org/downloads>  
(msi for windows)
- Create a new folder *c:\mongodb*, and install mongodb to the folder
- Create a new folder *data* inside the folder *mongodb* to store data

# Lab : Start MongoDB

- Open a cmd, goto folder `c:\mongodb\bin`
- Start monfoDB by typing  
`mongod.exe --dbpath c:\mongodb\data`
- Open another cmd, goto folder `c:\mongodb\bin` ,  
and type `mongo.exe` to open a management window.




```
c:\mongodb\bin>mongo
MongoDB shell version: 2.6.3
connecting to: test
> -
```

# Lab: Add a User to MongoDB

- Go to the management window and type:

```
> use course
> db.user.insert({name:"tom",pwd:"1234"})
> db.user.find({name:"tom"})
```



```
C:\mongodb\bin\mongo.exe

> use course
switched to db course
> db.user.insert({name:"tom",pwd:"1234"})
WriteResult({ "nInserted" : 1 })
> db.user.find({name:"tom"})
{ "_id" : ObjectId("53b8d1d514fbe94d1d3a8553"), "name" : "tom", "pwd" : "1234" }

>
```

# Lab: Close MongoDB

- Go to the management window and type:

```
> use admin  
> db.shutdownServer()
```

# **REST**

# What is REST?

- REST stands for *Representational State Transfer*.
- It describes an architecture for distributed information systems
- First described in the 2000 doctoral dissertation “Architectural Styles and the Design of Network-based Software Architectures” by [Roy Fielding](#).
- It’s a description of how the Web works and why it works well

# REST is about Architecture

- Software architecture
  - An abstraction of the elements, configurations, constraints, principles, and guidelines that govern a system's design and evolution.
- Software architectural style
  - A set of constraints that restrict a software architecture.

# RESTful Examples

- Public services with RESTful APIs:
  - Twitter, Netflix, Dropbox, Flickr, Amazon S3, ...
- Products or tools with RESTful APIs
  - Glassfish Application Server Admin, Selenium WebDriver, ...
- RESTful Frameworks
  - Jersey (JAX-RS), Restlet, Restify, APEX RESTful Services, ...

# Example: Amazon S3

90

- Amazon S3 is a Simple network Storage Service.
  - Provides cheap storage with read/write and pub/private access for the internet apps
- Support REST and SOAP APIs
  - REST API is quickly understood in 7-page tutorial
  - SOAP API WSDL is 7 pages of XML with no explanations at all!

# Why REST?

- Scalable
- Human and machine usable
- Language agnostic
- Globally accessible resources
- Intuitively understandable URIs, resources, and actions.


# REST Properties

- Client/server model – request/response
- Stateless: no memory of prior communications
- Cacheable
- Layered system: may use intermediary servers
- Code on demand (optional)
- Uniform interface
  - Request response style operations on named resources through self descriptive representations where state changes are via hyperlinks

# Uniform Interface


## The REST Triangle:

- Resources
- Methods
- Representations



# Uniform Interfaces - Resources

- Key abstract concept
- Identified by a URI
- Distinct from underlying storage
- Semantics fixed
- Value may change over time
- Can have multiple URIs
- Can have multiple representations
- Examples:
  - <http://example.org/NewOrleans/traffic/10>
  - <http://example.org/traffic/NewOrleans/I10>
  - <http://foo.com/store/orders>




# Resource Modeling - URIs

- Human readable (not necessary but it helps)
- Tends to form a hierarchy
- Use the query part appropriately
  - Use to search, filter, or possibly specify a mode
  - Identification of the resource is better in the path
    - (preferred) `http://example.com/orders/100234`
    - `http://example.com/orders?id=100234`
- Don't make them verbs!
  - (bad) `http://example.com/accounts/addaccount`

# Uniform Interface - Methods

- HTTP method Applying to the resource
  - **GET** retrieve
  - **PUT** update (or create)
  - **DELETE** delete
  - **POST** create sub resource
- Response codes: HTTP
  - 1xx, 2xx, 3xx, 4xx, 5xx



# REST Methods

- HTTP methods POST, GET, PUT, and DELETE are often compared with Create, Read, Update, and Delete (CRUD) operations associated with database functions:

REST	CRUD	SQL
GET	Read	SELECT
POST	Create	INSERT
PUT	Update or Create	UPDATE or INSERT
DELETE	Delete	DELETE

# A Example BBS System

- What message actions are required?
  - Create message
  - View message
  - View message for update
  - Delete message

# RESTful Methods

- How do the message actions map in REST?

Action	HTTP Method	URL
Create Message	POST	http://bbs/messages
View Message	GET	http://bbs/messages/12
Update Message	PUT	http://bbs/messages/12
Delete Message	DELETE	http://bbs/messages/12

# Design REST Methods

Method	Requirement
GET	Retrieve a resource. No modification should be done. No side effects allowed. Keep it safe.
POST	<i>Create or update a resource.</i> For non-safe, non-repeatable changes
PUT	<i>Update a resource.</i> Keep it repeatable with same results (idempotent).
DELETE	<i>Remove a resource.</i> Keep it repeatable with same results (idempotent)

# Idempotence

“Idempotent operations are operations that can be applied multiple times without changing the result”

- Idempotent tasks can be retried
  - Set Salary to 60K
    - If run twice, salary will still be set to 60k
- Non-Idempotent tasks cannot safely be retried
  - Retrieve current salary and increase by 10k
    - If run twice, salary will be increased 20k, not 10k

# Common Pattern for Methods

messages /

- GET - Retrieves list of all messages.
- POST - Create a new message .

messages/{mesno}/

- GET - Retrieves details for a specific message .
- PUT - Updates the specific message .
- DELETE - Deletes the message .


# Are They REST ?

- How do these actions map to HTTP and Web pages?

Action	HTTP Method	URL
Create Message	POST	http://bbs/msg.php?m=create
View Message	GET	http://bbs/msg.php?m=view&id=12
Update Message	POST	http://bbs/msg.php?m=update&Id=12
Delete Message	GET	http://bbs/msg.php?m=delete&id=12

# Uniform Interface - Representations

- Self-descriptive
- media type (Content-Type)
  - text/html
  - application/json
- Includes metadata
- Understood by all components
- May be for humans, machines or both
- Negotiated



# Representation Example

*Representation based on request type*

- Request for XML output

```
GET /resources/forte HTTP/1.1  
Host: example.com  
Accept: application/xml
```

- Request for HTML output

```
GET /resources/forte HTTP/1.1  
Host: example.com  
Accept: text/html
```

- Could be any format you wish to support.

# State Transfer Example

- Resources link together. State of the current system can be transferred along the links.

```
<resource self='http://example.com/resources/forte'>
  <name>forte</amount>
  <hostname>forte.example.com</hostname>
  <status>up</status>
  <account
    ref='http://example.com/accounts/1212'>1212</account>
    <user ref='http://example.com/profiles/jdoe'>jdoe</user>
</resource>
```

# REST vs. Web Services (RPC)

	Web Services	REST
Protocols	SOAP (Simple Object Access Protocol) over HTTP.	HTTP
Request Mechanism	XML over HTTP, usually POST	HTTP
Actions (verbs)	Many different actions, which are hidden within the request body.	Standard HTTP methods (GET, PUT, POST, DELETE)
Security	Additional SOAP-specific security layer.	Web server security
Web Server	HTTP and Web server are simple conduits with much of their power and capability are bypassed	HTTP and Web server exploited to fullest extent

# Example of RPC Application

- RPC Operations

```
getUser()    addUser()    removeUser() updateUser()  
getLocation() addLocation() removeLocation()  
updateLocation()   listUsers()  listLocations()  
findLocation()     findUser()
```

- RPC Client Code

```
exampleAppObject = new ExampleApp('example.com:1234')  
exampleAppObject.removeUser('001')
```

# Example of REST Application

- REST Define Resources

`http://example.com/users/  
http://example.com/users/{user} (1 for each user)  
http://example.com/findUserForm  
http://example.com/locations/  
http://example.com/locations/{location}  
http://example.com/findLocationForm`

- REST Client Code

```
userR = new Resource('http://example.com/users/001')  
userR.delete()
```

# Designing RESTful Services

- Key Principles
  - Everything gets a unique URI
  - Link resources together
  - Use standard methods
  - Resources have multiple representations
  - Communicate statelessly

# Access RESTful Services

- REST is meant for interoperability.
- Built on same technology as the internet.
- Can be consumed in many different ways
  - Web browser
  - Web services
  - Lightweight clients
  - Command line (curl, wget)

# REST Summary


- REST is an architectural style that
  - Focuses on resources
  - Minimizes verbs
  - Leverages existing Web server capabilities
- REST encourages
  - Loose coupling
    - Separation of representation and process
    - Separation of security and logging from application
  - Maximum exploitation of Web servers and HTTP

# **DEVELOP WEB APPS FOR CLOUD**


# Characteristics of Web Apps

- Today's focus on large-scale and ubiquitously useable Web Applications
  - Many Users – many languages – many cultures
  - Different access mechanisms
  - Many User Agents
- Presents large volume of interrelated information (including different media) and processes
  - Appropriate presentation
  - Progression through activities – finish one thing before starting another
- Growing and increasing complexity
  - Many product iterations/versions/refinements (calls for Reuse)
  - Many developers and operators, complex handling of temporal media (e.g. publishing of company news)
  - Customization, personalization, security issues

# Range of Complexity




# Problems with Large Web-based Projects



(Source: Epner, M., Cutter Consortium)

# Web Application Development



- Still Ad-Hoc instead of a disciplined procedure
  - Lacks rigor, systematic approach
- Lack between design model and Implementation model
- Design concepts get lost in the underlying model
- Short lifecycle of a web application -> maintenance and evolution issues -> reuse issues

# Web App Design


- Any complex system must begin with a cohesive fundamental design
- Multiple aspects of webapp design
  - Product Design: how the overall application will function, including:
  - Interaction (User-Interface) Design: The operation, look, and feel of the user interface
  - Software Design: The required software components and how they will interact

# Design Patterns

- A software **design pattern** is an abstraction of some commonly used software design
- This abstraction includes enough information to enable reuse of the pattern with new applications
- A pattern for architectural (high-level) design is also called an **architectural style**

# Model-View-Controller


- Many webapps are based on the well-known **model-view-controller** design pattern



- The MVC pattern specifies the responsibilities of three major components, and the nature of their interactions

# MVC Component Responsibilities

- View: Present the user interface
- Model: Maintain the application state
- Controller: Handle user actions



# Model Role

- The Model maintains the application state, which includes:
  - persistent information stored in databases
  - current information related to active sessions
- Responsibilities include:
  - applying rules / transactions to modify state
  - providing information to the View as required
  - taking instructions from the Controller

# View Role

- The View presents a user interface, including information and transactional controls
- Responsibilities include:
  - Present required information to user
  - Request data from Model as needed (to create displays for the user)


# Controller Role

- The Controller handles user actions
- Responsibilities include:
  - Handle user actions
  - Validate user requests (correctness, completeness)
  - Invoke Model components to handle requested transactions
  - Set the appropriate next View perspective

# Single-Page Web App

- One host (main) HTML file
- One screen at any given time
- Generating (or toggling) other screens dynamically, using JS
- Instead of navigating between different pages, one navigates between different sections of one web page
- Maintaining state – what's currently in the DOM
- Maintaining data – what's needed for generating DOM

# Building Web Applications




- **Host page:** A **HTML** page that we'd like to make interactive.
- **CSS** is used to decorate the main page.

# Building Web Applications (2)

- **Client-side script**: A JavaScript program that
  - registers handlers for relevant **DOM** events, such as inputs or mouse clicks
  - requests additional data from the server using **Ajax**;
  - integrates the responses with the web page using the **DOM**
  - Send back to the server the user's data using **Ajax**
- **Server side**: A program that supplies and receives the data
- Client-side script runs in browser, server-side using Node.js


# Web Development Cycles




# Security Requirements

- Privacy - All user information are protected
- Authentication/Access Control- Only authorized users are allowed to access the resources
- Integrity - User and application data cannot be tempered with
- Auditing - Keeping audit logs and audit trails and ensuring their integrity


# Applications Fit for Cloud




- On & off workloads (e.g. batch job)
- Over provisioned capacity is wasted
- Time to market can be cumbersome



- Successful services needs to grow/scale
- Keeping up w/ growth is big IT challenge
- Cannot provision hardware fast enough



- Unexpected/unplanned peak in demand
- Sudden spike impacts performance
- Can't over provision for extreme cases



- Services with micro seasonality trends
- Peaks due to periodic increased demand
- IT complexity and wasted capacity


# Cloud Apps Development

- Cloud software development
  - Design and develop an application for the cloud
  - Especially for the public PaaS cloud platforms
- Typical steps in cloud software development
  - Choose a development stack of technologies
  - Choose a cloud platform + services
  - Design the application for the cloud
  - Develop the application using the cloud APIs
  - Deploy and run the application in the cloud

# Transition to Cloud Development

- Transition to cloud development
  - New architecture (based on SOA)
  - New programming paradigms
    - E.g. NoSQL databases
  - New APIs
    - E.g. Amazon S3
  - New deployment model
    - Git + vendor-specific continuous integration process

# Moving an App to Cloud



# Example: WorkTitle

- A team cooperation tool
- Web-based cross-platforms, visit from browsers
- Single page application. Interface and feeling of native apps, such as drag and drop
- Instant message, every modification in one client will be automatically refreshed in other clients.
- Provide stable service
- Already more than 100 thousands users


# SPA of WorkTitle

The screenshot displays the Worktile Single Page Application (SPA) interface. The top navigation bar includes the 'Worktile' logo, a search bar, and user profile information for 'Jaden'. Below the header, a breadcrumb navigation shows the current location as '产品管理'. The main content area features a Kanban-style board with four columns:

- 产品设计**: Contains tasks like 'UI设计优化' (due 01-06), '收费系统、策略' (due 08-31), '换肤和背景' (due 10-30), '设计事件产品原型' (due 10-08), and '任务归档'.
- 通用**: Contains tasks like '活动时间流' (due 10-27), 'Dashboard优化' (due 10-25), '后台API编写' (due 10-18), 'Business Meeting' (due 10-15), and '样式优化' (due 10-02).
- 任务模块**: Contains tasks like '拖动成员分配任务' (due 10-17), '日历, 包括全局日历和项目日历, 任务在日历的创建' (due 10-01), '安全性' (due 10-15), '通过邮件创建任务', and '任务筛选'.
- 移动端**: Contains tasks like 'Andriod端' (due 11-30), 'iPad端' (due 12-31), 'Andriod Pad端', and a '新建任务' button.


On the right side, there is a vertical sidebar with a navigation menu. The '任务' section is currently selected, showing icons for '任务' (checked), '日历', '文件', '话题', '文档', and '简报'. Other sections like '产品', '日历', '文件', '话题', '文档', and '简报' are also listed but not selected.

# Front-side: AngularJs



- Two-way data binding
- Declarative semantic tags
- Modules
- Dependency injection

# Back-End of WorkTitle




- Preview service: Microsoft Office Web App
- Push service: ejabberd (core technology )

# Long Polling vs Short Polling vs Websocket

- Short polling : uses a timer and asks the server if new data is available
- Long Polling: uses an event based algorithm, more realtime.
- Extensible Messaging and Presence Protocol (**XMPP**) is a near-real-time communications protocol for message exchange.
- **ejabberd** is the world's most popular XMPP server. It can

# WorkTitle on Aliyun



# **Cloud Architecture Techniques**

**OpenStack/Docker**

# What is Virtualization?


- Virtualization
  - Running several virtual machines (virtual computers) inside a single physical machine
    - Supported by special software called *hypervisor*
  - Uses resources more efficiently
    - E.g. 12 GB physical RAM is shared to 6 virtual machines with 4 GB shared RAM each
    - Most applications use 1-5% of the CPU
      - A single shared CPU can serve thousands of users
  - Reduces costs due to better utilization

# Virtualization Definitions

- **Virtualization** is the ability to run multiple operating systems on a single physical system and share the underlying hardware resources
- Virtualization is used to improve IT throughput and costs by using physical resources as a pool from which virtual resources can be allocated.
- Virtual machine (VM) software is a program that emulates a physical machine
- A VM needs to act *exactly* like its physical machine
- A VM instance is simply a file that represents an actual machine and its state

# Virtualization Architecture

- A Virtual machine (VM) is an isolated runtime environment (guest OS and applications)
- Multiple virtual systems (VMs) can run on a single physical system



# OpenStack & Docker

Open Frameworks



PaaS  
Application Centric



IaaS  
Resource  
Centric



# OPENSTACK

# OpenStack

- An open source cloud computing framework
- It is primarily used to enable **infrastructure as a service**
- OpenStack began in 2010 as a joint project of Rackspace Hosting and NASA. Currently, it is managed by the OpenStack Foundation

# OpenStack Mission


To produce the  
**ubiquitous open source cloud computing  
platform**

that will meet the needs of public and private  
clouds regardless of size, by being simple to  
implement and massively scalable.


# Trends and Themes of OpenStack

- OpenStack Powered Planet – delivering on cloud interoperability and federation across vendors and geographies
- Applications on OpenStack – new community app catalog and first app guide provide more support for developers targeting OpenStack clouds
- Platform for the Next 10 Years – OpenStack gives users the ability to integrate and test new technologies like containers as they emerge

# OpenStack is Complicated



# But, Basic of OpenStack




# VMs as OpenStack Cloud Resources

- VMs run on top of hypervisors, and Agent process runs on physical machine beside hypervisor.
- User requests VM (resource) from controller, controller finds available resources in pool, contacts appropriate Agent
- Agent interprets command, instructs hypervisor to boot VM with user provided OS image
- Openstack returns handle to the user

# Components of OpenStack

- Compute (Nova)
- Object Storage (Swift)
- Block Storage (Cinder)
- Networking (Neutron)
- Dashboard (Horizon)
- Identity Service (Keystone)
- Telemetry (Ceilometer)
- Orchestration (Heat)
- Database (Trove)
- Bare Metal Provisioning (Ironic)
- Multiple Tenant Cloud Messaging (Zaqar)
- Elastic Map Reduce (Sahara)

# OpenStack Main Components



# Compute (Nova)


- Provision and manages virtual machines. It is a major part of an IaaS system.
  - **Nova-api service**: Accepts and responds to end user compute API calls.
  - **Nova-compute service**: A worker daemon that creates and terminates virtual machine instances through hypervisor APIs.
  - **Nova-scheduler**: Takes a virtual machine instance request from the queue and determines on which compute server host it runs.
  - **Rabbit MQ Server**: A central hub for passing messages between daemons

# Storage (Swift)

- Object Storage:
  - Object Storage provides a fully distributed, API-accessible storage platform that can be integrated directly into applications or used for backup, archiving and data retention.
- Block Storage:
  - Block Storage allows block of storages to be exposed and connected to compute instances for expanded storage, better performance and integration with enterprise storage platforms

# Networking(Neuron)

- Neutron is an OpenStack project to provide "networking as a service" between interface devices.
- It accepts and routes API requests to the appropriate OpenStack Networking plug-in devices for action.



# Dashboard (Horizon)

- Dashboard provides a web based user interface to OpenStack services including Nova, Swift, Keystone, etc.
- The OpenStack dashboard provides administrators and users a graphical interface to access, provision and automate cloud-based resources.

# Sample Dashboard

The screenshot shows the OpenStack Dashboard interface. On the left, there's a sidebar with navigation links: Project (selected), CURRENT PROJECT stackato, Manage Compute (Overview, Instances selected), Volumes, Images & Snapshots, Access & Security, Object Store (Containers). The main content area is titled 'Instances' and shows a table of running instances:


<input type="checkbox"/>	Instance Name	IP Address	Size	Keypair	Status	Task	Power State	Actions
<input type="checkbox"/>	stackato-266.data	10.0.3.2 208.75.128.146	m1.medium   4GB RAM   2 VCPU   40GB Disk	-	Active	None	Running	<button>Create Snapshot</button>
<input type="checkbox"/>	stackato-266.dean	10.0.2.18 208.75.128.145	m1.medium   4GB RAM   2 VCPU   40GB Disk	-	Active	None	Running	<button>Create Snapshot</button>
<input type="checkbox"/>	stackato-266.dean	10.0.5.14 208.75.128.144	m1.medium   4GB RAM   2 VCPU   40GB Disk	-	Active	None	Running	<button>Create Snapshot</button>
<input type="checkbox"/>	stackato-266.dean	10.0.4.10 208.75.128.143	m1.medium   4GB RAM   2 VCPU   40GB Disk	-	Active	None	Running	<button>Create Snapshot</button>
<input type="checkbox"/>	stackato-266.vrm1	10.0.3.6 208.75.128.142	m1.medium   4GB RAM   2 VCPU   40GB Disk	-	Active	None	Running	<button>Create Snapshot</button>

A success message box is visible in the top right corner: "Success: IP address 208.75.128.146 associated." Below the table, it says "Displaying 5 items".

# Identity Service (Keystone)

- Keystone is the identity service used by OpenStack for user authentication.
- It performs the following functions:
  - Tracking users and their permissions.
  - Providing a catalog of available services with their API endpoints.

# Keystone and Dashboard



# Image Service (Glance)

- The OpenStack Image Service provides discovery, registration and delivery services for disk and server images.
  - glance-api: Accepts Image API calls for image discovery, retrieval, and storage.
  - glance-registry: Stores, processes, and retrieves metadata about images.
  - Storage repository for image files: Various repository types are supported.


# Telemetry (Ceilometer)

- The OpenStack Telemetry service aggregates usage and performance data across the services deployed in an OpenStack cloud.

# Orchestration (Heat)

- OpenStack Heat program is to create a human- and machine-accessible service for managing the entire lifecycle of infrastructure and applications within OpenStack clouds.

# Architecture of OpenStack



# Use Cases of Openstack:


- IaaS compute platform
- Processing big data with tools like Hadoop
- Scaling compute up and down to meet demand for web resources and applications
- High-performance computing (HPC) environments processing diverse and intensive workloads
- Different users of OpenStack are AT&T, HP cloud services, PayPal, Rackspace, Sony online gaming systems, Yahoo, Wikimedia labs, Intel, NASA etc...

# Private Cloud of




- Availability: 99.99%, 2 breakdowns in 16 months
- Utilization rate: 40%(physical ). 1 PM : 12 VM
- The cloud spreads across 4 machine rooms. It has more 2000 VMs, 4500 physical cores (E5-2640), 50 T memory, 1200T virtual disk, 480 T block storage, 120 T object storage.

# Architecture of MI cloud



# Details of the Architecture



# More Facts about MI Cloud

- Compute node: DELL \_R720
  - CPU: E5-2640v2\*2(32 cores)
  - MEM: 16G\*24
  - Disk: 2\*600G SAS (Raid1) + 6\*4T (Raid5) SATA
  - Net Card: 1G \* 2 + 10G\*2 (Intel 82599EB 10-Gigabit SFI/SFP+ )
- Control node: DELL\_R620
  - CPU: E5-2630v2\*2 (24cores)
  - MEM: 16G\*4
  - Disk: 2\*600G SAS (Raid1) + 2\*240G SSD(Raid1)
  - Net Card: 1G \* 2 + 10G\*2 (Intel 82599EB 10-Gigabit SFI/SFP+ )
- OS : CentOS + RDO

# Walmart is powered by 140,000+ OpenStack Cores



"Last Thanksgiving, we did 1.5 billion page views... November and December are the two months where we have to satisfy all the needs of all the customers, and that's what we did – we built about 140,000 cores"

-Amandeep Singh Juneja, WalmartLabs

# PayPal' s transactions powered by OpenStack!



Havana  
> 12k hypervisors  
> 300k cores  
10+ availability zones  
15+ virtual private clouds  
> 1.6 pb block storage  
100% KVM  
100% OVS

Spring 2015

During the holidays for 2014, we were running 66% of PayPal front and mid-tier on OpenStack. That number is now 100%. **If you have doubts about OpenStack running in production, this is the proof point for you to say that OpenStack is doing financial transactions.” -Subbu Allamaraju, eBay Inc**



# DOCKER

# Docker


*Docker is an open-source project that automates the deployment of applications inside software containers, by providing an additional layer of abstraction and automation of operating system-level virtualization on Linux.*

[Source: [en.wikipedia.org](https://en.wikipedia.org)]


# What is Docker?

- Open source platform of light weight container technology for virtualization
- Docker allows you to package an application with all of its dependencies into a standardized unit for software development
- Applications:
  - Are dockerized and
  - Run on Docker Containers
    - from laptops to production servers on cloud
  - Their images are shared on Docker Hub
  - Apps can be linked (node -> mongo)

# Docker vs. Virtual Machine



Virtual Machine



Container

# Docker vs. Virtual Machine

- VMs require apps to be contained within the guest OS
  - Potentially 1:1 on VM:application
- Docker sits between the app and the host OS
  - Takes place of the hypervisor
  - Runs images as an isolated process
  - Only the app and its libraries are compartmentalized instead of requiring the guest OS share that space as well

# Why Linux Containers (LXC)


- Fast
  - Runtime performance near bare metal speeds
  - Management operations (run, stop , start, etc.) in seconds / milliseconds
- Flexible
  - Containerize a “system”
  - Containerize “application(s)”
- Lightweight
  - Just enough Operating System (JeOS)
  - Minimal per container penalty
- Inexpensive
  - Open source – free – lower TCO
  - Supported with out-of-the-box modern Linux kernel
- Ecosystem
  - Growing in popularity
  - Vibrant community & numerous 3<sup>rd</sup> party apps

# Containers Change the Game




# Docker Technology

- libvirt: Platform Virtualization
- LXC (LinuX Containers):  
Multiple isolated Linux  
systems (containers) on a  
single host
- Layered File System



# Docker Components



# Docker Parts

- Docker Engine
  - Client-server application
  - Handles running images, registries and containers
- Docker Hub
  - Public Docker registry
  - Holds a collection of existing images to be used publicly
  - Image examples: Ubuntu, Node.js, MongoDB, WordPress

# Docker Hub

- Public repository of Docker images
  - <https://hub.docker.com/>
- Integrates with Github and Bitbucket
  - Operates similar to them as well
- Hosts images that can be used and contributed to by the Docker community
  - Location that Docker pulls and pushes to

# Dockerfile

- Docker can build images automatically by reading the instructions from a Dockerfile.
- A Dockerfile is a text document that contains all the commands you would normally execute manually in order to build a Docker image.
- By calling docker build from your terminal, you can have Docker build your image step by step, executing the instructions successively

# Example Dockerfile

```
# Dockerfile to create a node image
FROM google/nodejs
MAINTAINER Carrotzpc
# Add files to the image
RUN mkdir -p /opt/nodejs
ADD . /opt/nodejs
WORKDIR /opt/nodejs
# Install the dependencies modules
RUN npm install
# Run make build
RUN make build
# Expose environment variables
ENV MONGO_CARROT_ADDR **LinkMe**
ENV MONGO_CARROT_PORT **LinkMe**
ENV MONGO_CARROT_DATABASE admin
ENV MONGO_CARROT_USER **ChangeMe**
ENV MONGO_CARROT_PASS **ChangeMe**
# Expose the container port
EXPOSE 5000
ENTRYPOINT ["node", "app.js"]
```


# Runing Platforms

- Various Linux distributions (Ubuntu, Fedora, RHEL, Centos, openSUSE, ...)
- Cloud (Amazon EC2, Google Compute Engine, Rackspace)
- Microsoft plans to integrate Docker with next release of Windows Server

# Running a Web App in Docker

- Similar to any other \*nix command line application
- Docker's example runs a Python Flask app:
  - \$ sudo docker run -d -P training/webapp python app.py
    - -P: map required network ports inside the container to the host
  - Check results with sudo docker ps -l:
    - CONTAINER ID IMAGE COMMAND CREATED STATUS PORTS NAMES  
bc533791f3f5 training/webapp:latest python app.py 5 seconds ago Up 2 seconds 0.0.0.0:49155->5000/tcp nostalgic\_morse
    - -l: tells docker ps to return the last container started

# Development Scenario



# MeanStack Use case



Open-Source Full-Stack Solution for MEAN  
Applications


# MeanStack

- Mean.js provides
  - Code generator to generate Mean App
  - Mean.js apps typically have
    - Node Js Server
    - Mongo DB database
  - Provides Dockerfile and fig.yml to run the app in Docker Containers
    - One Docker container for Node Js Server
    - One Docker container for Mongo DB Database

# MeanStack Configuration

```
fig.yml x  
  
web:  
  build: .  
  links:  
    - db  
  ports:  
    - "3000:3000"  
  environment:  
    NODE_ENV: development  
  
db:  
  image: mongo  
  ports:  
    - "27017:27017"
```

[fig.yml](#)



# Who uses Docker?

## Docker PAAS Providers



Microsoft Azure

**Google** Cloud Platform



dotCloud



StackDock



And many more...

# Who uses Docker?

As an Infrastructure Tool along side



CHEF™



Jenkins

# Docker Integration

- OpenStack + Docker
- Hadoop + Docker
- Spark + Docker
- Mesos + Docker
- .....


# OpenStack Deployment With Docker

- **Nova:** add a new nova docker driver.  
<https://github.com/stackforge/nova-docker>
- **HEAT:**
  - Add a new HEAT Resource:
  - HEAT docker driver interacts with docker server
- **Container as a Service**
  - A new service for OpenStack for manage docker container


Solutions are still under development

# Havana: nova-docker driver


- Nova driver which integrates with docker REST API on backend
- Glance translator to integrate docker images with Glance



# Icehouse: heat plugin for docker



# HEAT Docker Driver



- <https://github.com/MarouenMechtri/Docker-containers-deployment-with-OpenStack-Heat>
- <http://techs.enovance.com/7104/multi-tenant-docker-with-openstack-heat>